

# VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



**LAB REPORT**  
**on**

## **BIG DATA ANALYTICS** **(20CS6PEBDA)**

*Submitted by*

**Praveen Kumar S (1BM20CS413)**

*in partial fulfillment for the award of the degree of*  
**BACHELOR OF ENGINEERING**

*in*  
**COMPUTER SCIENCE AND ENGINEERING**



**B.M.S. COLLEGE OF ENGINEERING**

(Autonomous Institution under VTU)

**BENGALURU-560019**

**May-2022 to July-2022**

**B. M. S. College of Engineering,**  
**Bull Temple Road, Bangalore 560019**  
(Affiliated To Visvesvaraya Technological University, Belgaum)  
**Department of Computer Science and Engineering**



**CERTIFICATE**

This is to certify that the Lab work entitled “**BIG DATA ANALYTICS**” carried out by **Praveen Kumar S (1BM20CS413)** who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2022. The Lab report has been approved as it satisfies the academic requirements in respect of a **BIG DATA ANALYTICS** work prescribed for the said degree.

**Mrs. Antara Roy Choudhury**  
Designation  
Assistant Professor  
BMSCE, Bengaluru

**Dr. Jyothi S Nayak**  
Professor and Head  
Department of CSE  
BMSCE, Bengaluru

### Index Sheet

Sl. No.	Experiment Title	Page No.
1	Mongo DB	
2	Cassandra - Employee	
3	Cassandra - Library	
4	Hadoop Installation SS	
5	HDFS Commands	
6	Map Reduce for Temperature Mean and Avg	
7	Map Reduce for Top N Words	
8	Map Reduce to Demonstrate Join	
9	Word Count on Scala	
10	RDD and Flat Map for word Count	

### Course Outcome

CO1	Apply the concept of NoSQL, Hadoop or Spark for a given task
CO2	Analyze the Big Data and obtain insight using data analytics mechanisms.
CO3	Design and implement Big data applications by applying NoSQL, Hadoop or Spark

## WORKING WITH MONGODB

### I.CREATE DATABASE IN MONGODB.

**use myDB;**

Confirm the existence of your database

```
test>
>>> use myDB;
switched to db myDB
myDB>
>>>
```

**db;**

To list all databases

**show dbs;**

```
>>> show dbs;
admin      102 kB
config     12.3 kB
local      73.7 kB
myDB>
>>>
```

### I.CRUD (CREATE, READ, UPDATE, DELETE) OPERATIONS

1. To create a collection by the name “Student”. Let us take a look at the collection list prior to the creation of the new collection “Student”.

**db.createCollection(“Student”);**    =>    *sql equivalent*    **CREATE TABLE STUDENT(...);**

```
>>> db.createCollection("Student");
{ ok: 1 }
myDB>
>>>
```

1. To drop a collection by the name “Student”.

**db.Student.drop();**

1. Create a collection by the name “Students” and store the following data in it.

```
db.Student.insert({_id:1,StudName:"MichelleJacintha",Grade:"VII",Hobbies:"InternetSurfing"});
```

```
>>> db.Student.insertOne({ _id : 1, StudentName : "Bruce Wayne", Grade : "7", Hobbies : "Training"});  
{ acknowledged: true, insertedId: 1 }
```

1. Insert the document for “AryanDavid” in to the Students collection only if it does not already exist in the collection. However, if it is already present in the collection, then update the document with new values. (Update his Hobbies from “Skating” to “Chess”. ) Use “Update else insert” (if there is an existing document, it will attempt to update it, if there is no existing document then it will insert it).

```
db.Student.update({_id:3,StudName:"AryanDavid",Grade:"VII"},{$set:{Hobbies:"Skating"}},{upsert:true});
```

```
>>> db.Student.find();  
[  
  {  
    _id: 1,  
    StudentName: 'Bruce Wayne',  
    Grade: '7',  
    Hobbies: 'Training'  
  },  
  { _id: 2, StudentName: 'Clark Kent', Grade: '7', Hobbies: 'Skating' }  
]
```

```
>>> db.Student.updateOne({_id : 2, StudentName : "Clark Kent", Grade : "7"},{$set : {Hobbies : "Chess"}},{upset : true});  
{  
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 1,  
  modifiedCount: 1,  
  upsertedCount: 0  
}
```

```
>>> db.Student.find();  
[  
  {  
    _id: 1,  
    StudentName: 'Bruce Wayne',  
    Grade: '7',  
    Hobbies: 'Training'  
  },  
  { _id: 2, StudentName: 'Clark Kent', Grade: '7', Hobbies: 'Chess' }  
]
```

## 1. FIND METHOD

A. To search for documents from the “Students” collection based on certain search criteria.

```
db.Student.find({StudName:"Aryan David"});  
({cond..},{columns.. column:1, columnname:0} )
```

```
myDB>  
>>> db.Student.find({StudentName : "Bruce Wayne"});  
[  
  {  
    _id: 1,  
    StudentName: 'Bruce Wayne',  
    Grade: '7',  
    Hobbies: 'Training'  
  }  
]
```

B. To display only the StudName and Grade from all the documents of the Students collection. The identifier\_id should be suppressed and NOT displayed.

```
db.Student.find({}, {StudentName:1, Grade:1, _id:0});
```

```
myDB>  
>>> db.Student.find({}, {StudentName : 1, Grade : 1, _id :0});  
[  
  { StudentName: 'Bruce Wayne', Grade: '7' },  
  { StudentName: 'Clark Kent', Grade: '7' }  
]  
myDB>
```

C. To find those documents where the Grade is set to ‘VII’

```
db.Student.find({Grade:{$eq:'VII'}}).pretty();
```

```
myDB>
>>> db.Student.find({Grade : {$eq : "7"}});
[
  {
    _id: 1,
    StudentName: 'Bruce Wayne',
    Grade: '7',
    Hobbies: 'Training'
  },
  { _id: 2, StudentName: 'Clark Kent', Grade: '7', Hobbies: 'Chess' }
]
myDB>
```

D. To find those documents from the Students collection where the Hobbies is set to either 'Chess' or is set to 'Skating'.

```
db.Student.find({Hobbies : { $in: ['Chess','Skating']}}).pretty ();
```

```
myDB>
>>> db.Student.find({Hobbies : {$in : ["Chess","Skating"] }});
[ { _id: 2, StudentName: 'Clark Kent', Grade: '7', Hobbies: 'Chess' } ]
myDB>
```

E. To find documents from the Students collection where the StudName begins with "M".

```
db.Student.find({StudName:/^M/}).pretty();
```

```
myDB>
>>> db.Student.find({StudentName: /^B/});
[
  {
    _id: 1,
    StudentName: 'Bruce Wayne',
    Grade: '7',
    Hobbies: 'Training'
  }
]
```

F. To find documents from the Students collection where the StudName has an "e" in any position.

```
db.Student.find({StudName:/e/}).pretty();
```

```
myDB>
>>> db.Student.find({StudentName: /e/});
[
  {
    _id: 1,
    StudentName: 'Bruce Wayne',
    Grade: '7',
    Hobbies: 'Training'
  },
  { _id: 2, StudentName: 'Clark Kent', Grade: '7', Hobbies: 'Chess' }
]
myDB>
```

G. To find the number of documents in the Students collection.

```
db.Student.count();
```

```
myDB>
>>> db.Student.countDocuments();
2
myDB>
```

H. To sort the documents from the Students collection in the descending order of StudName.

```
db.Student.find().sort({StudName:-1}).pretty();
```

```
myDB>
>>> db.Student.find().sort({StudentName: -1});
[
  { _id: 2, StudentName: 'Clark Kent', Grade: '7', Hobbies: 'Chess' },
  {
    _id: 1,
    StudentName: 'Bruce Wayne',
    Grade: '7',
    Hobbies: 'Training'
  }
]
myDB>
```

### I.Import data from a CSV file

Given a CSV file “sample.txt” in the D:drive, import the file into the MongoDB collection, “SampleJSON”. The collection is in the database “test”.



```
mongoimport --db Student --collection airlines --type csv --headerline --file
/home/hduser/Desktop/airline.csv
```

### I.Export data to a CSV file

This command used at the command prompt exports MongoDB JSON documents from “Customers” collection in the “test” database into a CSV file “Output.txt” in the D:drive.

```
mongoexport --host localhost --db Student --collection airlines --csv --out
/home/hduser/Desktop/output.txt --fields “Year”, “Quarter”
```

### I.Save Method :

Save() method will insert a new document, if the document with the \_id does not exist. If it exists it will replace the existing document.

```
db.Students.save({StudName:”Vamsi”, Grade:”VI”})
```

### I. Add a new field to existing Document:

```
db.Students.update({_id:4},{ $set: {Location:”Network”}})
```

```
myDB>
>>> db.Student.update({_id : 1},{ $set : {Location : "Gotham City"}});
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

```
myDB>
>>> db.Student.find({_id:{$eq: 1}});
[
  {
    _id: 1,
    StudentName: 'Bruce Wayne',
    Grade: '7',
    Hobbies: 'Training',
    Location: 'Gotham City'
  }
]
myDB>
```

## I.Remove the field in an existing Document

```
db.Students.update({_id:4},{Sunset:{Location:"Network"}})
```

```
myDB>
>>> db.Student.update({_id : 1},{Sunset : {Location : "Gotham City"}});
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
myDB>
```

## I. Finding Document based on search criteria suppressing few fields

```
db.Student.find({_id:1},{StudName:1,Grade:1,_id:0});
```

```
myDB>
>>> db.Student.find({_id : 1}, {StudentName : 1, Grade : 1, _id : 0});
[ { StudentName: 'Bruce Wayne', Grade: '7' } ]
myDB>
```

To find those documents where the Grade is not set to 'VII'

```
db.Student.find({Grade:{$ne:'VII'}}).pretty();
```

```
>>> db.Student.find({Grade : {$ne : "7"}});
[
  {
    _id: ObjectId("6277c3e0bd8f013c5c3f84d1"),
    StudentName: 'Diana Prince',
    Grade: '8'
  }
]
myDB>
```

To find documents from the Students collection where the StudName ends with s.

```
db.Student.find({StudName:/s$/}).pretty();
```

```

myDB>
>>> db.Student.find({StudentName: /e$/});
[
  {
    _id: 1,
    StudentName: 'Bruce Wayne',
    Grade: '7',
    Hobbies: 'Training'
  },
  {
    _id: ObjectId("6277c3e0bd8f013c5c3f84d1"),
    StudentName: 'Diana Prince',
    Grade: '8'
  }
]
myDB>

```

**I.to set a particular field value to NULL**

```
db.Students.update({_id:3},{ $set:{Location:null}})
```

```

>>> db.Student.updateOne({_id : 1}, {$set : {Location : null}});
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
myDB>
>>> db.Student.find();
[
  {
    _id: 1,
    StudentName: 'Bruce Wayne',
    Grade: '7',
    Hobbies: 'Training',
    Location: null
  },

```

## I.Count the number of documents in Student Collections

```
db.Students.count()
```

```
>>> db.Student.count() ;  
3
```

## I.Count the number of documents in Student Collections with grade :VII

```
db.Students.count({Grade:"VII"})
```

```
myDB>  
>>> db.Student.count({Grade: "7"});  
1
```

retrieve first 3 documents

```
db.Students.find({Grade:"VII"}).limit(3).pretty();
```

```
>>> db.Student.find({Grade: "7"}).limit(3);  
[ { _id: 1, StudentName: 'Bruce Wayne', Grade: '7' } ]  
myDB>
```

Sort the document in Ascending order

```
db.Students.find().sort({StudName:1}).pretty();
```

```
myDB>  
>>> db.Student.find().sort({StudentName:1});  
[  
  { _id: 1, StudentName: 'Bruce Wayne', Grade: '7' },  
  { _id: 2, StudentName: 'Clark Kent', Grade: '9' },  
  { _id: 3, StudentName: 'Diana Prince', Grade: '10' }  
]  
myDB>
```

Note:

for desending order : db.Students.find().sort({StudName:-1}).pretty();

```
myDB>
>>> db.Student.find().sort({StudentName:-1});
[
  { _id: 3, StudentName: 'Diana Prince', Grade: '10' },
  { _id: 2, StudentName: 'Clark Kent', Grade: '9' },
  { _id: 1, StudentName: 'Bruce Wayne', Grade: '7' }
]
```

to Skip the 1<sup>st</sup> two documents from the Students Collections

```
db.Students.find().skip(2).pretty()
```

```
>>> db.Student.find().skip(2);
[ { _id: 3, StudentName: 'Diana Prince', Grade: '10' } ]
myDB>
```

XII. Create a collection by name “food” and add to each document add a “fruits” array

```
db.food.insert( { _id:1, fruits:['grapes','mango','apple'] } )
db.food.insert( { _id:2, fruits:['grapes','mango','cherry'] } )
db.food.insert( { _id:3, fruits:['banana','mango'] } )
```

```
>>> db.createCollection("food");
{ ok: 1 }
test>
>>> db.food.insertOne({_id : 1, fruits : ["Apple","Mango","Jack Fruit"]}));
{ acknowledged: true, insertedId: 1 }
test>
>>> db.food.insertOne({_id : 2, fruits : ["Cherry","Orange","Butter Fruit"]}));
{ acknowledged: true, insertedId: 2 }
test>
>>> db.food.insertOne({_id : 3, fruits : ["Banana","Water Melon"]}));
{ acknowledged: true, insertedId: 3 }
test>
>>>
```

To find those documents from the “food” collection which has the “fruits array” constitute of “grapes”, “mango” and “apple”.

```
db.food.find ( { fruits: ['grapes','mango','apple'] } ). pretty().
```

```
test>
>>> db.food.find({fruits:["Banana","Water Melon"]});
[ { _id: 3, fruits: [ 'Banana', 'Water Melon' ] } ]
test>
>>>
```

To find in “fruits” array having “mango” in the first index position.  
 db.food.find ( { 'fruits.1': 'grapes' } )

```
test>
>>> db.food.find({ 'fruits.0' : 'Banana' });
[ { _id: 3, fruits: [ 'Banana', 'Water Melon' ] } ]
test>
>>>
```

To find those documents from the “food” collection where the size of the array is two.

db.food.find ( { “fruits”: { \$size: 2 } } )

To find the document with a particular id and display the first two elements from the array “fruits”

db.food.find({\_id:1},{“fruits”:{ \$slice:2}})

```
test>
>>> db.food.find({ 'fruits' : { $size : 2 } });
[ { _id: 3, fruits: [ 'Banana', 'Water Melon' ] } ]
test>
>>>
```

To find all the documents from the food collection which have elements mango and grapes in the array “fruits”

db.food.find({fruits: { \$all: [ “mango”, “grapes” ] } })

```
test>
>>> db.food.find({fruits: { $all: [ "Cherry", "Orange" ] } }) ;
[ { _id: 2, fruits: [ 'Cherry', 'Orange', 'Butter Fruit' ] } ]
test>
>>>
```

update on Array:

using particular id replace the element present in the 1<sup>st</sup> index position of the fruits array with apple

```
db.food.update({_id:3},{ $set:{'fruits.1':'apple'}})
```

```
test>
>>> db.food.update({_id : 3}, {$set : {"fruits.1" : "Green Apple"}});
DeprecationWarning: Collection.update() is deprecated. Use updateOne,
updateMany, or bulkWrite.
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
test>
```

insert new key value pairs in the fruits array

```
db.food.update({_id:2},{ $push:{price:{grapes:80,mango:200,cherry:100}}})
```

```
test>
>>> db.food.update({_id : 3}, {$push : {price : {Banana : 20, GreenApple: 200}}});
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
test>
>>> db.food.find();
[
  { _id: 1, fruits: [ 'Apple', 'Mango', 'Jack Fruit' ] },
  { _id: 2, fruits: [ 'Cherry', 'Orange', 'Butter Fruit' ] },
  {
    _id: 3,
    fruits: [ 'Banana', 'Green Apple' ],
    price: [ {}, { Banana: 20, GreenApple: 200 } ]
  }
]
test>
```

Note: perform query operations using - pop, addToSet, pullAll and pull

## XII. Aggregate Function :

Create a collection Customers with fields custID, AcctBal, AcctType.  
Now group on “custID” and compute the sum of “AccBal”.

```
>>> db.Customer.find();
{ "_id" : ObjectId("629449502b957d283eee6404"), "CustId" : 1, "AcctBal" : 1000, "AcctType" : "Savings" }
{ "_id" : ObjectId("629449872b957d283eee6405"), "CustId" : 1, "AcctBal" : 2000, "AcctType" : "Current" }
{ "_id" : ObjectId("6294499e2b957d283eee6406"), "CustId" : 2, "AcctBal" : 50000, "AcctType" : "Current" }
{ "_id" : ObjectId("629449d12b957d283eee6407"), "CustId" : 2, "AcctBal" : 5000, "AcctType" : "Savings" }
>>>
```

```
db.Customers.aggregate ( { $group : { _id : "$custID", TotAccBal : { $sum : "$AccBal" } } } );
```

```
>>> db.Customer.aggregate({ $group : { _id : "$CustId", TotalAccBal :
{ $sum : "$AcctBal" } } });
{ "_id" : 2, "TotalAccBal" : 55000 }
{ "_id" : 1, "TotalAccBal" : 3000 }
>>>
```

match on AcctType:”S” then group on “CustID” and compute the sum of “AccBal”.

```
db.Customers.aggregate ( { $match: { AcctType: "S" } }, { $group : { _id : "$custID", TotAccBal :
{ $sum : "$AccBal" } } } );
```

```
>>> db.Customer.aggregate( { $match: { AcctType: "Savings" } }, { $group : { _id
: "$custID", TotalAccBal : { $sum : "$AcctBal" } } });
{ "_id" : null, "TotalAccBal" : 6000 }
```

match on AcctType:”S” then group on “CustID” and compute the sum of “AccBal” and  
total balance greater than 1200.

```
db.Customers.aggregate ( { $match: { AcctType: "S" } }, { $group : { _id : "$custID", TotAccBal :
{ $sum : "$AccBal" } } }, { $match: { TotAccBal : { $gt: 1200 } } } );
```

```
>>> db.Customer.aggregate( { $match: { AcctType: "Savings" } }, { $group : { _id
: "$custID", TotalAccBal : { $sum : "$AcctBal" } } }, { $match: { TotalAccBal :
{ $gt: 1200 } } } );
{ "_id" : null, "TotalAccBal" : 6000 }
>>>
```



# Cassandra Program - 1

## 1. Create a key space by name Employee

```
cqlsh> CREATE KEYSPACE Empyolees WITH REPLICATION = { 'class' : 'SimpleStrategy',  
'replication_factor' : 1 };
```

```
cqlsh> DESCRIBE KEYSPACES;
```

```
system_schema crud    project system_distributed system_traces  
system_auth  system student empyolees
```

```
cqlsh> USE Employees;
```

## 2. Create a column family by name Employee-Info with attributes Emp\_Id Primary Key, Emp\_Name, Designation, Date\_of\_Joining, Salary, Dept\_Name

```
cqlsh:employees> CREATE TABLE Employee_Info (  
    ... Emp_Id int PRIMARY KEY,  
    ... Emp_Name text,  
    ... Designation text,  
    ... Date_Of_Joining timestamp,  
    ... Salary int,  
    ... Dept_Name text  
    ... );
```

```
cqlsh:employees> DESCRIBE TABLES;
```

```
employee_info
```

```
cqlsh:employees> DESCRIBE TABLE Employee_Info;
```

```
CREATE TABLE employees.employee_info (  
    emp_id int PRIMARY KEY,  
    date_of_joining timestamp,  
    dept_name text,  
    designation text,  
    emp_name text,  
    salary int  
) WITH bloom_filter_fp_chance = 0.01  
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
    AND comment = "  
    AND compaction = {'class':  
'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32',  
'min_threshold': '4'}  
    AND compression = {'chunk_length_in_kb': '64', 'class':  
'org.apache.cassandra.io.compress.LZ4Compressor'}
```

```

AND crc_check_chance = 1.0
AND dclocal_read_repair_chance = 0.1
AND default_time_to_live = 0
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair_chance = 0.0
AND speculative_retry = '99PERCENTILE';

```

### 3. Insert the values into the table in batch

```

cqlsh:employees> BEGIN BATCH
... INSERT INTO Employee_Info
(Emp_Id,Emp_Name,Designation,Date_of_Joining,Salary,Dept_Name) VALUES (1,'Bruce
Wayne','CEO','2022-04-22',100000,'Management')
... INSERT INTO Employee_Info
(Emp_Id,Emp_Name,Designation,Date_of_Joining,Salary,Dept_Name) VALUES (2,'Clark
Kent','Senior Software Engineer','2022-04-24',70000,'Developemt')
... INSERT INTO Employee_Info
(Emp_Id,Emp_Name,Designation,Date_of_Joining,Salary,Dept_Name) VALUES (3,'Diana
Prince','Jr Software Engineer','2022-04-30',70000,'Developemt')
... INSERT INTO Employee_Info
(Emp_Id,Emp_Name,Designation,Date_of_Joining,Salary,Dept_Name) VALUES (4,'Aurthr
Curry','Senior Manager','2022-05-30',70000,'Developemt')
... APPLY BATCH;

```

```

cqlsh:employees> SELECT * FROM Employee_Info;

```

emp_id	date_of_joining	dept_name	designation	emp_name	salary
1	2022-04-21 18:30:00.000000+0000	Management	CEO	Bruce Wayne	100000
2	2022-04-23 18:30:00.000000+0000	Developemt	Senior Software Engineer	Clark Kent	70000
4	2022-05-29 18:30:00.000000+0000	Developemt	Senior Manager	Aurthr Curry	70000
121	2022-06-29 18:30:00.000000+0000	Accounts	Accountant	Barry Allen	60000
3	2022-04-29 18:30:00.000000+0000	Developemt	Jr Software Engineer	Diana Prince	70000

### 4. Update Employee name and Department of Emp-Id 121

```

cqlsh:employees> UPDATE Employee_Info SET Emp_Name = 'Wally West', dept_name = 'HR'
WHERE Emp_id = 121;

```

emp_id	date_of_joining	dept_name	designation	emp_name	salary
1	2022-04-21 18:30:00.000000+0000	Management	CEO	Bruce Wayne	100000
2	2022-04-23 18:30:00.000000+0000	Developemt	Senior Software Engineer	Clark Kent	70000
4	2022-05-29 18:30:00.000000+0000	Developemt	Senior Manager	Aurthr Curry	70000
121	2022-06-29 18:30:00.000000+0000	HR	Accountant	Wally West	60000
3	2022-04-29 18:30:00.000000+0000	Developemt	Jr Software Engineer	Diana Prince	70000

## 5. Sort the details of Employee records based on salary

```
cqlsh:employees> CREATE TABLE Employee_Info (  
    ... Emp_Id int,  
    ... Emp_Name text,  
    ... Designation text,  
    ... Date_Of_Joining timestamp,  
    ... Salary int,  
    ... Dept_Name text,  
    ... PRIMARY KEY (Emp_Id , Salary)  
    ... ) WITH CLUSTERING ORDER BY (Salary desc);
```

```
cqlsh:employee> select * from Employee_Info;
```

emp_id	date_of_joining	dept_name	designation	emp_name	salary
121	2022-06-29 18:30:00.000000+0000	HR	Accountant	Wally West	60000
3	2022-04-29 18:30:00.000000+0000	Development	Jr Software Manager	Diana Prince	70000
2	2022-04-23 18:30:00.000000+0000	Management	Senior Software Manager	Clark Kent	70000
4	2022-05-29 18:30:00.000000+0000	Development	Senior Manager	Aurthur Curry	70000
1	2022-04-21 18:30:00.000000+0000	Management	CEO	Bruce Wayne	100000

## 6. Alter the schema of the table Employee\_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.

```
cqlsh:employee> ALTER TABLE Employee_Info ADD Projects text;
```

```
cqlsh:employee> select * from Employee_Info;
```

emp_id	date_of_joining	dept_name	designation	emp_name	projects	salary
1	2022-04-21 18:30:00.000000+0000	Management	CEO	Bruce Wayne	null	100000
2	2022-04-23 18:30:00.000000+0000	Management	Senior Software Manager	Clark Kent	null	70000
4	2022-05-29 18:30:00.000000+0000	Development	Senior Manager	Aurthur Curry	null	70000
121	2022-06-29 18:30:00.000000+0000	HR	Accountant	Wally West	null	60000
3	2022-04-29 18:30:00.000000+0000	Development	Jr Software Manager	Diana Prince	null	70000

## 7. Update the altered table to add project names.

```
cqlsh:employee> UPDATE Employee_Info SET Projects='Research' WHERE Emp_id=1 and  
salary=100000.0;
```

```
cqlsh:employee> select * from Employee_Info;
```

```
cqlsh:employee> select * from Employee_Info;
```

emp_id	date_of_joining	dept_name	designation	emp_name	projects	salary
1	2022-04-21 18:30:00.000000+0000	Management	CEO	Bruce Wayne	Research	100000
2	2022-04-23 18:30:00.000000+0000	Management	Senior Software Manager	Clark Kent	null	70000
4	2022-05-29 18:30:00.000000+0000	Development	Senior Manager	Aurthur Curry	null	70000
121	2022-06-29 18:30:00.000000+0000	HR	Accountant	Wally West	null	60000
3	2022-04-29 18:30:00.000000+0000	Development	Jr Software Manager	Diana Prince	null	70000

```
cqlsh:employee> UPDATE Employee_Info SET Projects='Data Migration' WHERE Emp_id=2
and salary=70000.0;
```

```
cqlsh:employee> UPDATE Employee_Info SET Projects='Data analysis' WHERE Emp_id=3
and salary=70000.0;
```

```
cqlsh:employee> UPDATE Employee_Info SET Projects='Reporting' WHERE Emp_id=121 and
salary=60000.0;
```

```
cqlsh:employee> UPDATE Employee_Info SET Projects='Research' WHERE Emp_id=4 and
salary=70000.0;
```

```
cqlsh:employee> select * from Employee_Info;
```

emp_id	date_of_joining	dept_name	designation	emp_name	projects	salary
1	2022-04-21 18:30:00.000000+0000	Management	CEO	Bruce Wayne	Research	100000
2	2022-04-23 18:30:00.000000+0000	Management	Senior Software Manager	Clark Kent	Data Migration	70000
4	2022-05-29 18:30:00.000000+0000	Development	Senior Manager	Aurthur Curry	Data analysis	70000
121	2022-06-29 18:30:00.000000+0000	HR	Accountant	Wally West	Reporting	60000
3	2022-04-29 18:30:00.000000+0000	Development	Jr Software Manager	Diana Prince	Research	70000

## 8 Create a TTL of 15 seconds to display the values of Employees

```
cqlsh:employee> INSERT INTO Employee_Info(Emp_id, Emp_Name, Designation, Date_Of_Joining,
salary, Dept_name) VALUES (5,'John Jones','CTO','2022-04-01',80000.0,'Space Station') using ttl 15;
```

```
cqlsh:employee> select ttl(Emp_Name) from Employee_Info Where Emp_id=5;
```

```
ttl(emp_name)
```

```
-----
```

## Cassandra Program - 2

### 1 Create a key space by name Library

```
bmsce@bmsce-Precision-T1700:~$ Cassandra/apache-cassandra-3.11.0/bin
bash: Cassandra/apache-cassandra-3.11.0/bin: Is a directory
bmsce@bmsce-Precision-T1700:~$ Cassandra/apache-cassandra-3.11.0/bin/
bash: Cassandra/apache-cassandra-3.11.0/bin/: Is a directory
bmsce@bmsce-Precision-T1700:~$ cd Cassandra/apache-cassandra-3.11.0/bin/
bmsce@bmsce-Precision-T1700:~/Cassandra/apache-cassandra-3.11.0/bin$ ./cqlsh
Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.11.4 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cqlsh> create keyspace library with replication = {
... 'class': 'SimpleStrategy', 'replication_factor': 1
... };
cqlsh> describe keyspaces
```

```
system_schema system student system_traces
system_auth library system_distributed
```

### 2. Create a column family by name Library-Info with attributes Stud\_Id Primary Key, Counter\_value of type Counter, Stud\_Name, Book-Name, Book-Id, Date\_of\_issue

```
cqlsh:library> create table library_info(stud_id int, counter_value counter, stud_name text, book_name
text, book_id int, date_of_issue date, primary key(stud_id, stud_name, book_name, book_id,
date_of_issue));
```

```
cqlsh:library> describe library_info
```

```
CREATE TABLE library.library_info (
  stud_id int,
  stud_name text,
  book_name text,
  book_id int,
  date_of_issue date,
  counter_value counter,
  PRIMARY KEY (stud_id, stud_name, book_name, book_id, date_of_issue)
) WITH CLUSTERING ORDER BY (stud_name ASC, book_name ASC, book_id ASC, date_of_issue
ASC)
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy',
'max_threshold': '32', 'min_threshold': '4'}
AND compression = {'chunk_length_in_kb': '64', 'class':
'org.apache.cassandra.io.compress.LZ4Compressor'}
AND crc_check_chance = 1.0
AND dclocal_read_repair_chance = 0.1
```

```

AND default_time_to_live = 0
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair_chance = 0.0
AND speculative_retry = '99PERCENTILE';

```

### 3. Insert the values into the table in batch

### 4. Display the details of the table created and increase the value of the counter

```

cqlsh:library> update library_info set counter_value = counter_value + 1 where stud_id = 1 and
stud_name = 'Bruce' and book_name = 'Game of Thrones' and book_id = 1 and date_of_issue = '2022-04-
20';

```

```

cqlsh:library> select * from library_info;

```

stud_id	stud_name	book_name	book_id	date_of_issue	counter_value
1	Bruce	Game of Thrones	1	2022-04-20	1

(1 rows)

```

cqlsh:library> update library_info set counter_value = counter_value + 1 where stud_id = 2 and
stud_name = 'Clark' and book_name = 'Song of Ice and Fire' and book_id = 2 and date_of_issue = '2022-
04-21';

```

```

cqlsh:library> select * from library_info;

```

stud_id	stud_name	book_name	book_id	date_of_issue	counter_value
1	Bruce	Game of Thrones	1	2022-04-20	1
2	Clark	Song of Ice and Fire	2	2022-04-21	1

(2 rows)

```

cqlsh:library> update library_info set counter_value = counter_value + 1 where stud_id = 112 and
stud_name = 'Diana' and book_name = 'BDA' and book_id = 3 and date_of_issue = '2022-05-04';

```

```

cqlsh:library> select * from library_info;

```

stud_id	stud_name	book_name	book_id	date_of_issue	counter_value
1	Bruce	Game of Thrones	1	2022-04-20	1
2	Clark	Song of Ice and Fire	2	2022-04-21	1
112	Diana	BDA	3	2022-05-04	1

(3 rows)

## 5. Write a query to show that a student with id 112 has taken a book “BDA” 2 times.

```
cqlsh:library> update library_info set counter_value = counter_value + 1 where stud_id = 112 and  
stud_name = 'Diana' and book_name = 'BDA' and book_id = 3 and date_of_issue = '2022-05-04';  
cqlsh:library> select * from library_info;
```

stud_id	stud_name	book_name	book_id	date_of_issue	counter_value
1	Bruce	Game of Thrones	1	2022-04-20	1
2	Clark	Song of Ice and Fire	2	2022-04-21	1
112	Diana	BDA	3	2022-05-04	2

(3 rows)

```
cqlsh:library> select * from library_info where stud_id = 112;
```

stud_id	stud_name	book_name	book_id	date_of_issue	counter_value
112	Diana	BDA	3	2022-05-04	2

(1 rows)

## 6. Export the created column to a csv file

```
cqlsh:library> copy library_info (stud_id, stud_name, book_name, book_id, date_of_issue,  
counter_value) to '/home/bmsce/Desktop/data.csv';  
Using 11 child processes
```

Starting copy of library.library\_info with columns [stud\_id, stud\_name, book\_name, book\_id, date\_of\_issue, counter\_value].

Processed: 4 rows; Rate: 21 rows/s; Avg. rate: 21 rows/s

4 rows exported to 1 files in 0.200 seconds.

## 7. Import a given csv dataset from local file system into Cassandra column family

```
cqlsh:library> copy library_info (stud_id, stud_name, book_name, book_id, date_of_issue,  
counter_value) from '/home/bmsce/Desktop/data1.csv';  
Using 11 child processes
```

Starting copy of library.library\_info with columns [stud\_id, stud\_name, book\_name, book\_id, date\_of\_issue, counter\_value].

Processed: 4 rows; Rate: 7 rows/s; Avg. rate: 11 rows/s

4 rows imported from 1 files in 0.381 seconds (0 skipped).

```
cqlsh:library> select * from library_info;
```

stud_id	stud_name	book_name	book_id	date_of_issue	counter_value
1	Bruce	Game of Thrones	1	2022-04-20	1
2	Clark	Song of Ice and Fire	2	2022-04-21	1
112	Diana	BDA	3	2022-05-04	2
1	Bruce	Game of Thrones	1	2022-04-20	1
2	Clark	Song of Ice and Fire	2	2022-04-21	1
112	Diana	BDA	3	2022-05-04	2



#### 4. Screenshot of Hadoop installed

```

Administrator: C:\WINDOWS\system32\cmd.exe
C:\> cd C:\Program Files\Apache Software Foundation\hadoop
C:\Program Files\Apache Software Foundation\hadoop> java -Djava.library.path=C:\Program Files\Apache Software Foundation\hadoop\bin -jar hadoop-2.7.0-hadoop2\bin\hadoop.dll
2022-07-14 23:16:21,692 INFO server.AbstractConnector: Started ServerConnector@740abb5{HTTP/1.1, (http/1.1)}{0.0.0.0:8042}
2022-07-14 23:16:21,692 INFO server.Server: Started @20289ms
2022-07-14 23:16:21,692 INFO webapp.WebApps: Web app node started at 8042
2022-07-14 23:16:21,694 INFO nodemanager.NodeStatusUpdaterImpl: Node ID assigned is : Praveen:63150
2022-07-14 23:16:21,698 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2022-07-14 23:16:21,711 INFO client.DefaultNMHARMAFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8031
2022-07-14 23:16:22,023 INFO security.NMContainerTokenSecretManager: Rolling master-key for container-tokens, got key with id 495823401
2022-07-14 23:16:22,024 INFO security.NMTokenSecretManagerInNM: Rolling master-key for container-tokens, got key with id 490100392
2022-07-14 23:16:22,024 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as Praveen:63150 with total resource <memory:8192, vCores:8>

```

```
Administrator: C:\WINDOWS\system32\cmd.exe
Microsoft Windows [Version 10.0.22000.675]
(c) Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>cd \

C:\>cd hadoop-3.3.3\sbin

C:\hadoop-3.3.3\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\hadoop-3.3.3\sbin>jps
10260 NameNode
11124 Jps
3732 ResourceManager
7076 NodeManager
3756 DataNode

C:\hadoop-3.3.3\sbin>
```

```
Administrator: C:\WINDOWS\system32\cmd.exe

C:\hadoop-3.3.3\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\hadoop-3.3.3\sbin>jps
10260 NameNode
11124 Jps
3732 ResourceManager
7076 NodeManager
3756 DataNode

C:\hadoop-3.3.3\sbin>hdfs dfs -mkdir /test

C:\hadoop-3.3.3\sbin>hdfs dfs -ls /
Found 12 items
drwxr-xr-x - Praveen supergroup      0 2022-07-14 21:34 /JoinProIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 22:30 /JoinProOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:30 /TempAvg
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:32 /TempAvgOp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:48 /TempMeanIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:49 /TempMeanOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:15 /TonNOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:04 /TopNIp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:02 /WordCountIp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:08 /WordCountOp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 23:19 /test
drwx----- - Praveen supergroup      0 2022-07-12 08:05 /tmp

C:\hadoop-3.3.3\sbin>
```

```
Administrator: C:\WINDOWS\system32\cmd.exe

C:\hadoop-3.3.3\sbin>hdfs dfs -rmdir /test

C:\hadoop-3.3.3\sbin>hdfs dfs -ls /
Found 11 items
drwxr-xr-x - Praveen supergroup      0 2022-07-14 21:34 /JoinProIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 22:30 /JoinProOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:30 /TempAvg
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:32 /TempAvgOp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:48 /TempMeanIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:49 /TempMeanOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:15 /TonNOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:04 /TopNIp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:02 /WordCountIp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:08 /WordCountOp
drwx----- - Praveen supergroup      0 2022-07-12 08:05 /tmp

C:\hadoop-3.3.3\sbin>hdfs dfs -cat /WordCount/*
cat: `/WordCount/*': No such file or directory

C:\hadoop-3.3.3\sbin>hdfs dfs -cat /WordCountIp/*
cat: `/WordCountIp/*': No such file or directory

C:\hadoop-3.3.3\sbin>hdfs dfs -cat /WordCountIp/*
Hello I'm Batman
Hey Flash How are you?
How is your brother Barry Allen?
Hey SuperMan How are you?
How is Lois Lane?
C:\hadoop-3.3.3\sbin>
```

```
Administrator: C:\WINDOWS\system32\cmd.exe
C:\hadoop-3.3.3\sbin>hdfs dfs -cp /WordCoundIp /CopyWC

C:\hadoop-3.3.3\sbin>hdfs dfs -ls /
Found 12 items
drwxr-xr-x - Praveen supergroup      0 2022-07-14 23:22 /CopyWC
drwxr-xr-x - Praveen supergroup      0 2022-07-14 21:34 /JoinProIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 22:30 /JoinProOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:30 /TempAvg
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:32 /TempAvgOp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:48 /TempMeanIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:49 /TempMeanOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:15 /TonNOP
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:04 /TopNIP
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:02 /WordCoundIp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:08 /WordCountOp
drwx----- - Praveen supergroup      0 2022-07-12 08:05 /tmp

C:\hadoop-3.3.3\sbin>hdfs dfs -cat /CopyWC/*
Hello I'm Batman
Hey Flash How are you?
How is your brother Barry Allen?
Hey SuperMan How are you?
How is Lois Lane?
C:\hadoop-3.3.3\sbin>
```

```
Administrator: C:\WINDOWS\system32\cmd.exe
Hey Flash How are you?
How is your brother Barry Allen?
Hey SuperMan How are you?
How is Lois Lane?
C:\hadoop-3.3.3\sbin>hdfs dfs -cp /WordCoundIp /CopyWC

C:\hadoop-3.3.3\sbin>hdfs dfs -ls /
Found 12 items
drwxr-xr-x - Praveen supergroup      0 2022-07-14 23:22 /CopyWC
drwxr-xr-x - Praveen supergroup      0 2022-07-14 21:34 /JoinProIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 22:30 /JoinProOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:30 /TempAvg
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:32 /TempAvgOp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:48 /TempMeanIp
drwxr-xr-x - Praveen supergroup      0 2022-07-14 20:49 /TempMeanOp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:15 /TonNOP
drwxr-xr-x - Praveen supergroup      0 2022-07-12 09:04 /TopNIP
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:02 /WordCoundIp
drwxr-xr-x - Praveen supergroup      0 2022-07-12 08:08 /WordCountOp
drwx----- - Praveen supergroup      0 2022-07-12 08:05 /tmp

C:\hadoop-3.3.3\sbin>hdfs dfs -cat /CopyWC/*
Hello I'm Batman
Hey Flash How are you?
How is your brother Barry Allen?
Hey SuperMan How are you?
How is Lois Lane?
C:\hadoop-3.3.3\sbin>hdfs dfs -get /WordCoundIp/* C:/New

C:\hadoop-3.3.3\sbin>
```

## 6. From the following link extract the weather data

<https://github.com/tomwhite/hadoop-book/tree/master/input/ncdc/all>

Create a Map Reduce program to

**a) find average temperature for each year from NCDC data set.**

### **//AverageDriver.java**

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class AverageDriver {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Please Enter the input and output parameters");
            System.exit(-1);
        }
        Job job = new Job();
        job.setJarByClass(AverageDriver.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.setMapperClass(AverageMapper.class);
        job.setReducerClass(AverageReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

### **//AverageMapper.java**

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class AverageMapper extends Mapper<LongWritable, Text,
Text, IntWritable> {
    public static final int MISSING = 9999;
    public void map(LongWritable key, Text value,
```

```

        Mapper<LongWritable, Text, Text, IntWritable>.Context context)
            throws IOException, InterruptedException {
    int temperature;
    String line = value.toString();
    String year = line.substring(15, 19);
    if (line.charAt(87) == '+') {
        temperature = Integer.parseInt(line.substring(88, 92));
    } else {
        temperature = Integer.parseInt(line.substring(87, 92));
    }
    String quality = line.substring(92, 93);
    if (temperature != 9999 && quality.matches("[01459]"))
        context.write(new Text(year), new
            IntWritable(temperature));
    }
}

```

### **//AverageReducer.java**

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class AverageReducer extends Reducer<Text, IntWritable,
Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values,
        Reducer<Text, IntWritable, Text, IntWritable>.Context context)
            throws IOException, InterruptedException {
        int max_temp = 0;
        int count = 0;
        for (IntWritable value : values) {
            max_temp += value.get();
            count++;
        }
        context.write(key, new IntWritable(max_temp / count));
    }
}

```

```
Select Administrator: C:\WINDOWS\system32\cmd.exe

Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=57
CPU time spent (ms)=1185
Physical memory (bytes) snapshot=574767104
Virtual memory (bytes) snapshot=823296000
Total committed heap usage (bytes)=370147328
Peak Map Physical memory (bytes)=339447808
Peak Map Virtual memory (bytes)=465547264
Peak Reduce Physical memory (bytes)=235319296
Peak Reduce Virtual memory (bytes)=357748736

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=73867
File Output Format Counters
  Bytes Written=8

C:\hadoop-3.3.3\sbin>hadoop dfs -cat /TempAvgOp/*
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
1901  46

C:\hadoop-3.3.3\sbin>
```

**b) find the mean max temperature for every month**

### **//MeanMaxDriver.java**

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class MeanMaxDriver {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Please Enter the input and output parameters");
            System.exit(-1);
        }
        Job job = new Job();
        job.setJarByClass(MeanMaxDriver.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.setMapperClass(MeanMaxMapper.class);
        job.setReducerClass(MeanMaxReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

```
}
```

### **//MeanMaxMapper.java**

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class MeanMaxMapper extends Mapper<LongWritable, Text,
Text, IntWritable> {
    public static final int MISSING = 9999;
    public void map(LongWritable key, Text value,
                    Mapper<LongWritable, Text, Text, IntWritable>.Context context)
                    throws IOException, InterruptedException {
        int temperature;
        String line = value.toString();
        String month = line.substring(19, 21);
        if (line.charAt(87) == '+') {
            temperature = Integer.parseInt(line.substring(88, 92));
        } else {
            temperature = Integer.parseInt(line.substring(87, 92));
        }
        String quality = line.substring(92, 93);
        if (temperature != 9999 && quality.matches("[01459]"))
            context.write(new Text(month), new
                IntWritable(temperature));
    }
}
```

### **//MeanMaxReducer.java**

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class MeanMaxReducer extends Reducer<Text, IntWritable,
Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values,
                    Reducer<Text, IntWritable, Text, IntWritable>.Context context)
                    throws IOException, InterruptedException {
```

}

```
C:\hadoop-3.3.3\sbin>
```



**7. For a given Text file, create a Map Reduce program to sort the content in an alphabetic order listing only top 10 maximum occurrences of words.**

### **//Driver Code**

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

public class TopN {
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        String[] otherArgs = (new GenericOptionsParser(conf, args)).getRemainingArgs();
        if (otherArgs.length != 2) {
            System.err.println("Usage: TopN <in> <out>");
            System.exit(2);
        }
        Job job = Job.getInstance(conf);
        job.setJobName("Top N");
        job.setJarByClass(TopN.class);
        job.setMapperClass(TopNMapper.class);
        job.setReducerClass(TopNReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
        FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

### **//Reducer**

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
```

```

public class TopNMapper extends Mapper<Object, Text, Text, IntWritable> {
    private static final IntWritable one = new IntWritable(1);

    private Text word = new Text();

    private String tokens = "[_!$#<>\\^=\\[\\]\\|*\\/\\\\\\,;\\.\\|\\-:()?!\\\"'"]";

    public void map(Object key, Text value, Mapper<Object, Text, Text, IntWritable>.Context
context) throws IOException, InterruptedException {
        String cleanLine = value.toString().toLowerCase().replaceAll(this.tokens, " ");
        StringTokenizer itr = new StringTokenizer(cleanLine);
        while (itr.hasMoreTokens()) {
            this.word.set(itr.nextToken().trim());
            context.write(this.word, one);
        }
    }
}

```

## //Reducer

```

import java.io.IOException;
import java.util.HashMap;
import java.util.Map;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
import utils.MiscUtils;

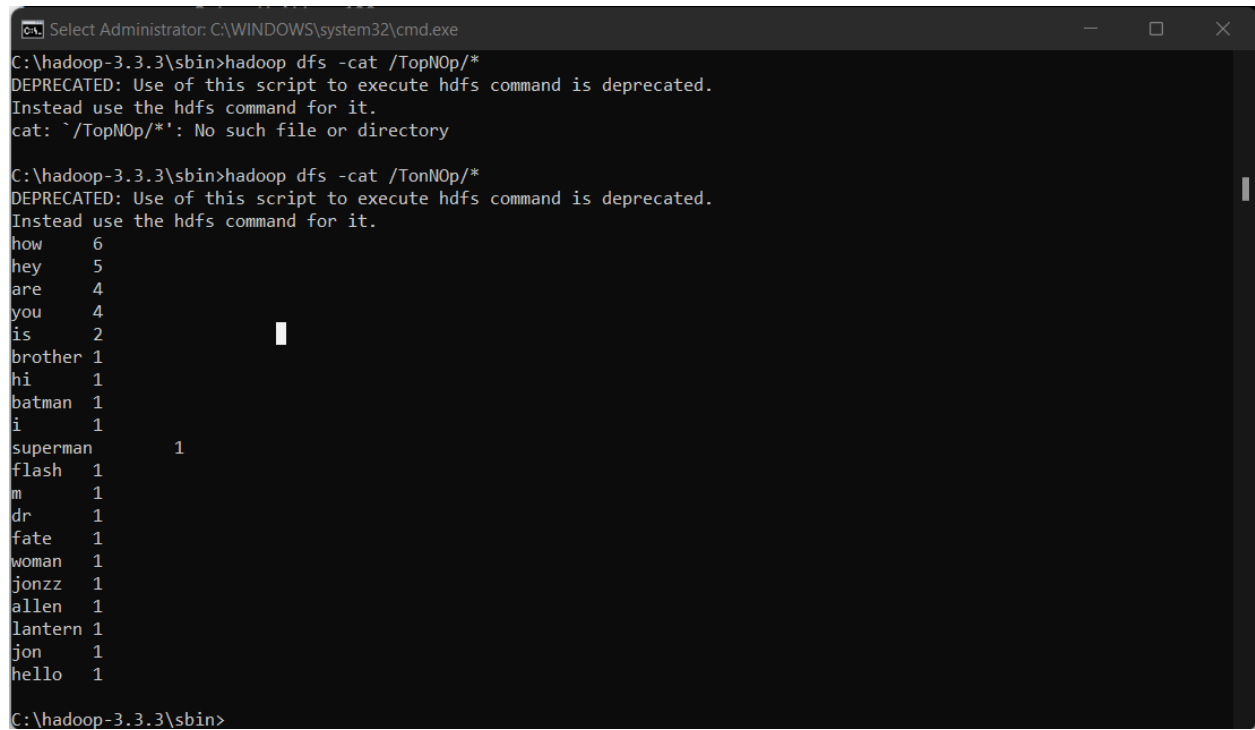
public class TopNReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    private Map<Text, IntWritable> countMap = new HashMap<>();

    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable, Text,
IntWritable>.Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values)
            sum += val.get();
        this.countMap.put(new Text(key), new IntWritable(sum));
    }

    protected void cleanup(Reducer<Text, IntWritable, Text, IntWritable>.Context context) throws
IOException, InterruptedException {
        Map<Text, IntWritable> sortedMap = MiscUtils.sortByValues(this.countMap);
        int counter = 0;
    }
}

```

```
for (Text key : sortedMap.keySet()) {  
    if (counter++ == 20)  
        break;  
    context.write(key, sortedMap.get(key));  
}  
}  
}
```



The screenshot shows a Windows command prompt window titled "Select Administrator: C:\WINDOWS\system32\cmd.exe". The user is in the directory C:\hadoop-3.3.3\sbin and runs the command `hadoop dfs -cat /TopNOp/*`. The output shows a deprecation warning and an error: `cat: ~/TopNOp/*: No such file or directory`. Then, the user runs `hadoop dfs -cat /TonOp/*`, which also shows a deprecation warning. The output then displays a list of words and their frequencies:

```
how 6  
hey 5  
are 4  
you 4  
is 2  
brother 1  
hi 1  
batman 1  
i 1  
superman 1  
flash 1  
m 1  
dr 1  
fate 1  
woman 1  
jonzz 1  
allen 1  
lantern 1  
jon 1  
hello 1
```

The prompt ends with `C:\hadoop-3.3.3\sbin>`.

## 8. Create a Map Reduce program to demonstrating join operation

### // JoinDriver.java

```
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.mapred.libMultipleInputs;
import org.apache.hadoop.util.*;

public class JoinDriver extends Configured implements Tool {
    public static class KeyPartitioner implements Partitioner < TextPair,
        Text > {
        @Override
        public void configure(JobConf job) {}
        @Override
        public int getPartition(TextPair key, Text value, int numPartitions) {
            return (key.getFirst().hashCode() & Integer.MAX_VALUE) %
                numPartitions;
        }
    }
    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 3) {
            System.out.println("Usage: <Department Emp Strength input> <
                Department Name input > < output > ");
            return -1;
        }
        JobConf conf = new JobConf(getConf(), getClass());
        conf.setJobName("Join 'Department Emp Strength input' with
            'Department Name
            input'");
        Path AInputPath = new Path(args[0]); Path BInputPath = new Path(args[1]); Path
        outputPath = new Path(args[2]); MultipleInputs.addInputPath(conf, AInputPath,
        TextInputFormat.class,
            Posts.class); MultipleInputs.addInputPath(conf, BInputPath,
        TextInputFormat.class,
            User.class); FileOutputFormat.setOutputPath(conf, outputPath);
        conf.setPartitionerClass(KeyPartitioner.class);
        conf.setOutputValueGroupingComparator(TextPair.FirstComparator.cl ass);
        conf.setMapOutputKeyClass(TextPair.class); conf.setReducerClass(JoinReducer.class);
        conf.setOutputKeyClass(Text.class); JobClient.runJob(conf);
        return 0;
    }
    public static void main(String[] args) throws Exception {
```

```

        int exitCode = ToolRunner.run(new JoinDriver(), args);
        System.exit(exitCode);
    }
}

```

## // JoinReducer.java

```

import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
public class JoinReducer extends MapReduceBase implements
Reducer < TextPair, Text, Text,
Text > {
    @Override
    public void reduce(TextPair key, Iterator < Text > values,
        OutputCollector < Text, Text >
        output, Reporter reporter)
        throws IOException {
        Text nodeId = new Text(values.next());
        while (values.hasNext()) {
            Text node = values.next();
            Text outValue = new Text(nodeId.toString() + "\t\t" + node.toString());
            output.collect(key.getFirst(), outValue);
        }
    }
}

```

## // JoinReducer.java

```

import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
public class JoinReducer extends MapReduceBase implements
Reducer < TextPair, Text, Text,
Text > {
    @Override
    public void reduce(TextPair key, Iterator < Text > values,
        OutputCollector < Text, Text >
        output, Reporter reporter)
        throws IOException {
        Text nodeId = new Text(values.next());
        while (values.hasNext()) {

```

```

        Text node = values.next();
        Text outValue = new Text(nodeId.toString() + "\t\t" + node.toString());
        output.collect(key.getFirst(), outValue);
    }
}
}

```

### **//Posts.java**

```

import java.io.IOException;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
public class Posts extends MapReduceBase implements
Mapper < LongWritable, Text, TextPair,
Text > {
    @Override
    public void map(LongWritable key, Text value,
        OutputCollector < TextPair, Text > output,
        Reporter reporter)
        throws IOException {
        String valueString = value.toString();
        String[] SingleNodeData = valueString.split("\t");
        output.collect(new TextPair(SingleNodeData[3], "0"), new Text(SingleNodeData[9]));
    }
}

```

### **// TextPair.java**

```

import java.io.*;
import org.apache.hadoop.io.*;
public class TextPair implements WritableComparable < TextPair > {
    private Text first;
    private Text second;
    public TextPair() {
        set(new Text(), new Text());
    }
    public TextPair(String first, String second) {
        set(new Text(first), new Text(second));
    }
    public TextPair(Text first, Text second) {
        set(first, second);
    }
    public void set(Text first, Text second) {
        this.first = first;
        this.second = second;
    }
}

```

```

    }
    public Text getFirst() {
        return first;
    }
    public Text getSecond() {
        return second;
    }
    @Override
    public void write(DataOutput out) throws IOException {
        first.write(out);
        second.write(out);
    }
    @Override
    public void readFields(DataInput in ) throws IOException {
        first.readFields( in );
        second.readFields( in );
    }
    @Override
    public int hashCode() {
        return first.hashCode() * 163 + second.hashCode();
    }
    @Override
    public boolean equals(Object o) {
        if (o instanceof TextPair) {
            TextPair tp = (TextPair) o;
            return first.equals(tp.first) && second.equals(tp.second);
        }
        return false;
    }
    @Override
    public String toString() {
        return first + "\t" + second;
    }
    @Override
    public int compareTo(TextPair tp) {
        int cmp = first.compareTo(tp.first);
        if (cmp != 0) {
            return cmp;
        }
        return second.compareTo(tp.second);
    }
}

```

## // TextPairComparator

```
public static class Comparator extends WritableComparator {
    private static final Text.Comparator TEXT_COMPARATOR = new
    Text.Comparator();
    public Comparator() {
        super(TextPair.class);
    }
    @Override
    public int compare(byte[] b1, int s1, int l1,
        byte[] b2, int s2, int l2) {
        try {
            int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1, s1);
            int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) + readVInt(b2, s2);
            int cmp = TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2,
                firstL2);
            if (cmp != 0) {
                return cmp;
            }
            return TEXT_COMPARATOR.compare(b1, s1 + firstL1, l1 - firstL1,
                b2, s2 + firstL2, l2 - firstL2);
        } catch (IOException e) {
            throw new IllegalArgumentException(e);
        }
    }
}

static {
    WritableComparator.define(TextPair.class, new Comparator());
}

public static class FirstComparator extends WritableComparator {
    private static final Text.Comparator TEXT_COMPARATOR = new
    Text.Comparator();
    public FirstComparator() {
        super(TextPair.class);
    }
    @Override
    public int compare(byte[] b1, int s1, int l1,
        byte[] b2, int s2, int l2) {
        try {
            int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1, s1);
            int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) + readVInt(b2, s2);
            return TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2, firstL2);
        } catch (IOException e) {
            throw new IllegalArgumentException(e);
        }
    }
}
```



```

    }
}
@Override
public int compare(WritableComparable a, WritableComparable b) {
    if (a instanceof TextPair && b instanceof TextPair) {
        return ((TextPair) a).first.compareTo(((TextPair) b).first);
    }
    return super.compare(a, b);
}
}
}

```

The screenshot shows a Windows command prompt window titled "Administrator: C:\WINDOWS\system32\cmd.exe". The output of a Hadoop job is displayed, including statistics for reduce output, memory usage, and shuffle errors. Below the job output, the user runs the command `C:\hadoop-3.3.3\sbin>hadoop dfs -cat /JoinProOp/*`, which results in a deprecation warning and a listing of files in the HDFS directory. The listing shows three files: A11 (100 bytes, Development), B12 (50 bytes, Testing), and C13 (30 bytes, HR). The prompt then shows the user running `C:\hadoop-3.3.3\sbin>`.

```

Administrator: C:\WINDOWS\system32\cmd.exe

    Reduce output records=4
    Spilled Records=16
    Shuffled Maps =4
    Failed Shuffles=0
    Merged Map outputs=4
    GC time elapsed (ms)=290
    CPU time spent (ms)=2352
    Physical memory (bytes) snapshot=1600139264
    Virtual memory (bytes) snapshot=2287669248
    Total committed heap usage (bytes)=1253572608
    Peak Map Physical memory (bytes)=370651136
    Peak Map Virtual memory (bytes)=490999808
    Peak Reduce Physical memory (bytes)=234283008
    Peak Reduce Virtual memory (bytes)=359739392

Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0

File Input Format Counters
    Bytes Read=0
File Output Format Counters
    Bytes Written=82

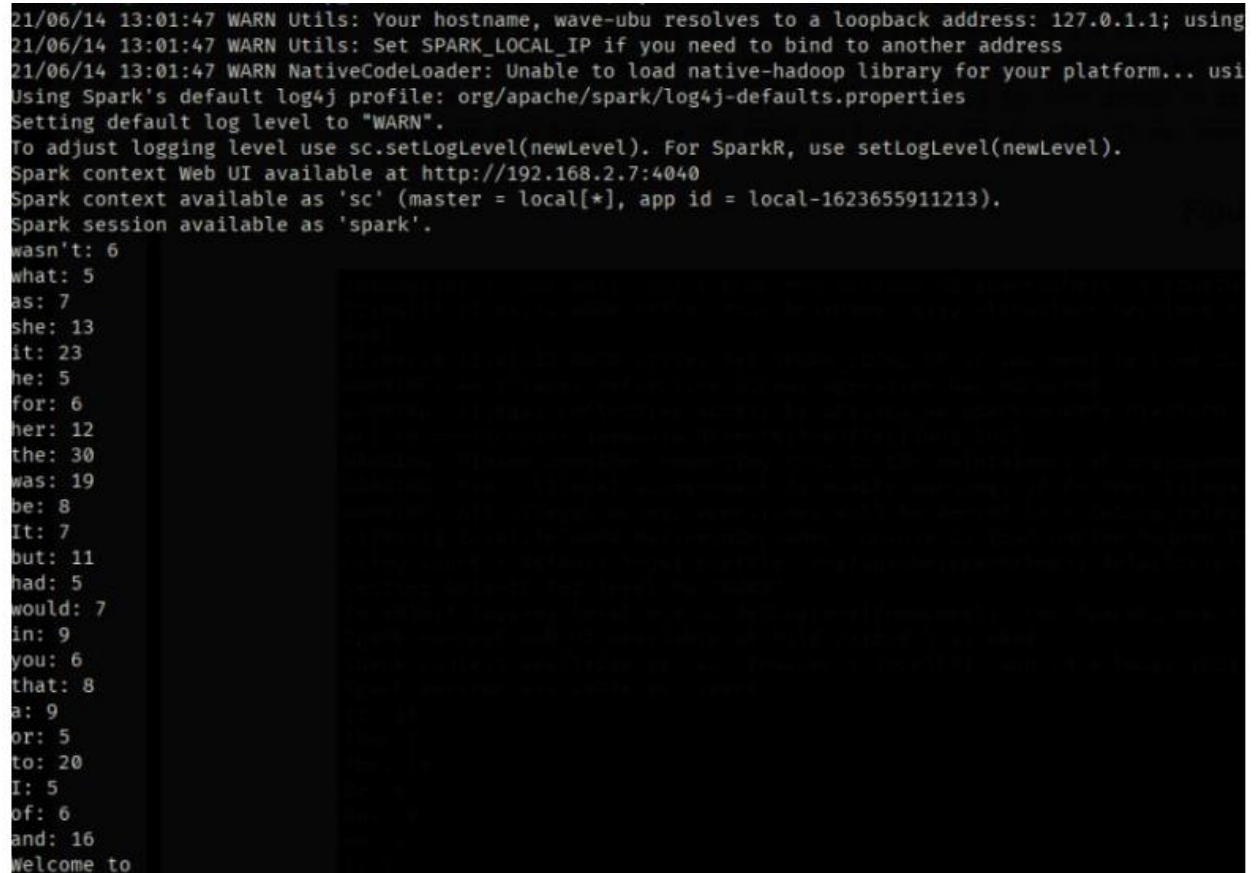
C:\hadoop-3.3.3\sbin>hadoop dfs -cat /JoinProOp/*
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
A11    100      Development
B12    50       Testing
C13    30       HR
Dept_ID Total_Employee  Dept_Name
C:\hadoop-3.3.3\sbin>

```

## 9. Program to print word count on scala shell and print “Hello world” on scala IDE

```
scala> println("Hello World!");  
Hello World!
```

```
val data=sc.textFile("Test.txt")  
data.collect;  
val splitdata = data.flatMap(line => line.split(" "));  
splitdata.collect;  
val mapdata = splitdata.map(word => (word,1));  
mapdata.collect;  
val reducedata = mapdata.reduceByKey(_+_);  
reducedata.collect;
```



The screenshot shows a terminal window with the following content:

```
21/06/14 13:01:47 WARN Utils: Your hostname, wave-ubu resolves to a loopback address: 127.0.1.1; using  
21/06/14 13:01:47 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address  
21/06/14 13:01:47 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... usi  
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
Spark context Web UI available at http://192.168.2.7:4040  
Spark context available as 'sc' (master = local[*], app id = local-1623655911213).  
Spark session available as 'spark'.  
wasn't: 6  
what: 5  
as: 7  
she: 13  
it: 23  
he: 5  
for: 6  
her: 12  
the: 30  
was: 19  
be: 8  
It: 7  
but: 11  
had: 5  
would: 7  
in: 9  
you: 6  
that: 8  
a: 9  
or: 5  
to: 20  
I: 5  
of: 6  
and: 16  
Welcome to
```

## 10. Using RDD and Flat Map count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark

```
val textFile = sc.textFile("/home/bhoom/Desktop/wc.txt")
val counts = textFile.flatMap(line => line.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)
import scala.collection.immutable.ListMap
val sorted = ListMap(counts.collect.sortWith(_._2 > _._2):_*) // sort in descending order based on values
println(sorted)
for((k,v) <- sorted)
{
  if(v > 4)
  {
    print(k+", ")
    print(v)
    println()
  }
}
```

```
scala> val textfile = sc.textFile("/home/sam/Desktop/abc.txt")
textfile: org.apache.spark.rdd.RDD[String] = /home/sam/Desktop/abc.txt MapPartitionsRDD[8] at textFile at <console>:25

scala> val counts = textfile.flatMap(line => line.split(" ")).map(word => (word,1)).reduceByKey(_+_ )
counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[11] at reduceByKey at <console>:26

scala> import scala.collection.immutable.ListMap
import scala.collection.immutable.ListMap

scala> val sorted = ListMap(counts.collect.sortWith(_._2>_._2):_*)
sorted: scala.collection.immutable.ListMap[String,Int] = ListMap(hello -> 3, apple -> 2, unicorn -> 1, world -> 1)

scala> println(sorted)
ListMap(hello -> 3, apple -> 2, unicorn -> 1, world -> 1)
```